

Gradient Descent, Regularization

Fitting Models as Optimization

- Loss: compare fitted to true value

$$\ell(y, \hat{y}) = (y - \hat{y})^2$$

Fitting Models as Optimization

- Loss: compare fitted to true value (w/ parameter)

$$\ell(y, X\beta) = (y - X\beta)^2$$

Fitting Models as Optimization

- Loss: compare fitted to true value (over data)

$$\sum_{i=1}^n \ell(y_i, x_i\beta) = \sum_{i=1}^n (y_i - x_i\beta)^2$$

Fitting Models as Optimization

- Optimization problem:

$$\hat{\beta} = \arg \min \sum_{i=1}^n (y_i - x_i \beta)^2$$


Gradient Descent

- Goal: minimize the loss
- Calculus approach to doing it:
 - Loss is a function of parameters
 - Take derivative (gradient) of that function (wrt parameters)
 - *Set equal to zero*
 - *Use algebra to solve it*


Gradient Descent

- Goal: minimize the loss
- (Approximate) calculus approach to doing it:
 - Loss is a function of parameters
 - Take derivative (gradient) of that function (wrt parameters)
 - *Derivative tells you where to go to decrease function*
 - *Follow the path to the solution*


When will Gradient Descent Work?



A




B




C


When will Gradient Descent Work?



A




B




C


When will Gradient Descent Work?



A





B



C

Gradient Descent

- Convex problems guaranteed to work
 - Can you roll a ball down the hill?
 - Will you end up in a unique spot if you do?



Do we only care about loss?

- What would happen if we are able to get the loss to zero on the training data? Is that good?

Do we only care about loss?

- Regularization lets us trade *within sample fit* and *complexity*.
 - Will a complex model generalize well?
 - Will a model that fits the training data perfectly generalize well?

Regularization: Complexity penalties

- Complexity penalties:

$$\hat{\beta} = \arg \min \sum_{i=1}^n (y_i - x_i \beta)^2 + \lambda \text{penalty}(\beta)$$

- What do we do to optimize this? When is this possible?

Regularization: Complexity penalties

$$\hat{\beta} = \arg \min \sum_{i=1}^n (y_i - x_i \beta)^2 + \lambda \text{penalty}(\beta)$$

“Ridge” $\text{penalty}(\beta) = \sum_j \beta_j^2$

“Lasso” $\text{penalty}(\beta) = \sum_j |\beta_j|$

Regularization: Complexity penalties

“Ridge” $\text{penalty}(\beta) = \sum_j \beta_j^2$

“Lasso” $\text{penalty}(\beta) = \sum_j |\beta_j|$

- How do these measure complexity?

Regularization: Complexity penalties

“L-Zero” $\text{penalty}(\beta) = \sum_j I(\beta_j \neq 0)$

- Why not use this penalty?

Other types of regularization

- Not going all the way in optimization
 - Stop before the optimum
 - Don't use all the data; e.g. Stochastic Gradient Descent (SGD)
 - Footnote on SGD: developed decades ago when machines were too weak to use all data, these days we find ourselves back in this situation with "Big Data"